

Compressed Sensing

Nicholas Chiapputo

EENG 5850.001

April 29, 2021

Motivation

- Conventional Image Capture:
 - Capture tons of data
 - ...throw most of it away
- Why not only capture the important information?
 - “Why go to so much effort to acquire **all** the data when **most** of what we get will be thrown away? Can’t we just **directly measure** the part that won’t end up being thrown away?” -- David Donoho
 - Problem: What data is important?
- Solution: Compressed Sensing
 - Compression during capture (on the fly)
 - Randomly select data to retrieve
 - Sparse optimization

Compressed Sensing

- Process
 - Take random projections (measurements) of the original signal
 - Recover the signal using optimization techniques
- How can we sample below the Nyquist rate?
 - Nyquist-Shannon theorem assumes a broadband signal
 - The signal is assumed to be “sparse” (most natural signals are) or compressible
 - Sparse: only a few frequencies are present
 - Transform into a domain in which it has a sparse representation
 - A signal is sparse if it most of its components are at or near zero
- Reconstruction
 - Sparse optimization

Problem

- Underdetermined linear system:

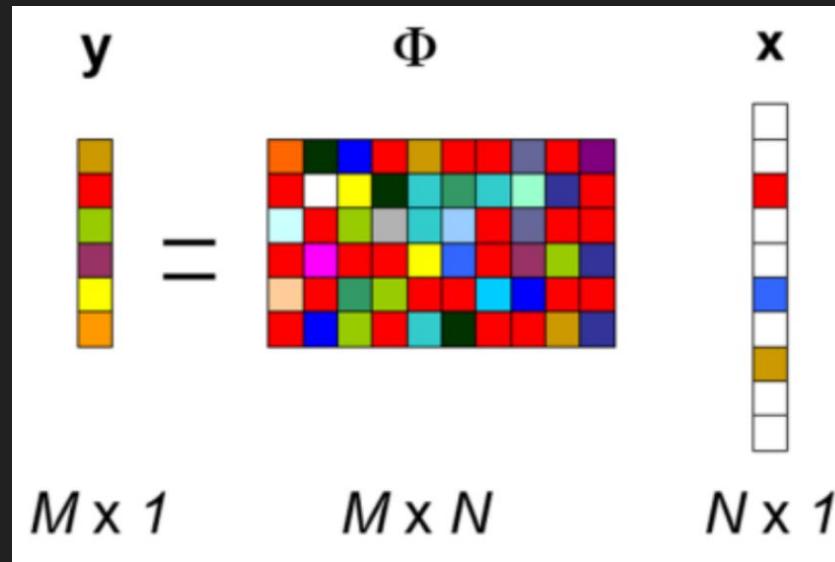
$$\begin{bmatrix} \boldsymbol{y} \end{bmatrix} = \begin{bmatrix} \boldsymbol{A} \end{bmatrix} \begin{bmatrix} \boldsymbol{x} \end{bmatrix}$$

Mx1 measurements MxN sensing matrix

- $k \ll M \ll N$ $N \times 1$ reconstructed components (sparse)
 - \mathbf{x} is sparse with k non-zero entries
 - \mathbf{A} is often also called Φ
 - Underdetermined -- infinite solutions
 - Projection of N -dimensional vector into an M dimensional space

Problem

- Underdetermined linear system:

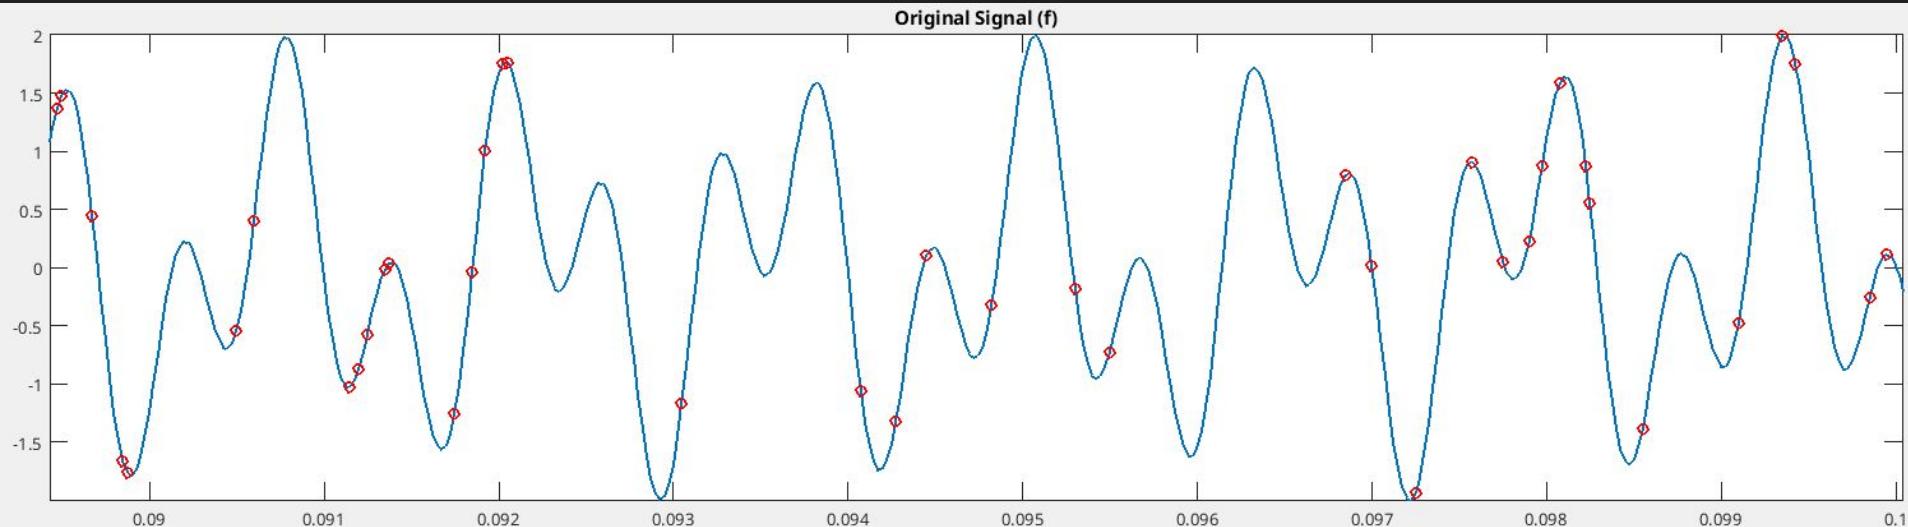


Visualized problem from [2]

How do we reconstruct?

- Solve the underdetermined equation $\mathbf{y} = \Phi\mathbf{x}$
- Ideal optimization: minimize $\|\mathbf{x}\|_0$ (number of non-zero elements in \mathbf{x})
 - This is an NP-hard problem!
- Other Options:
 - $\|\mathbf{x}\|_1$: L₁-norm -- least absolute deviations or Basis Pursuit (sparse solutions)
 - $\|\mathbf{x}\|_2$: L₂-norm -- least squares (fastest)
- Method is known as *Basis Pursuit*

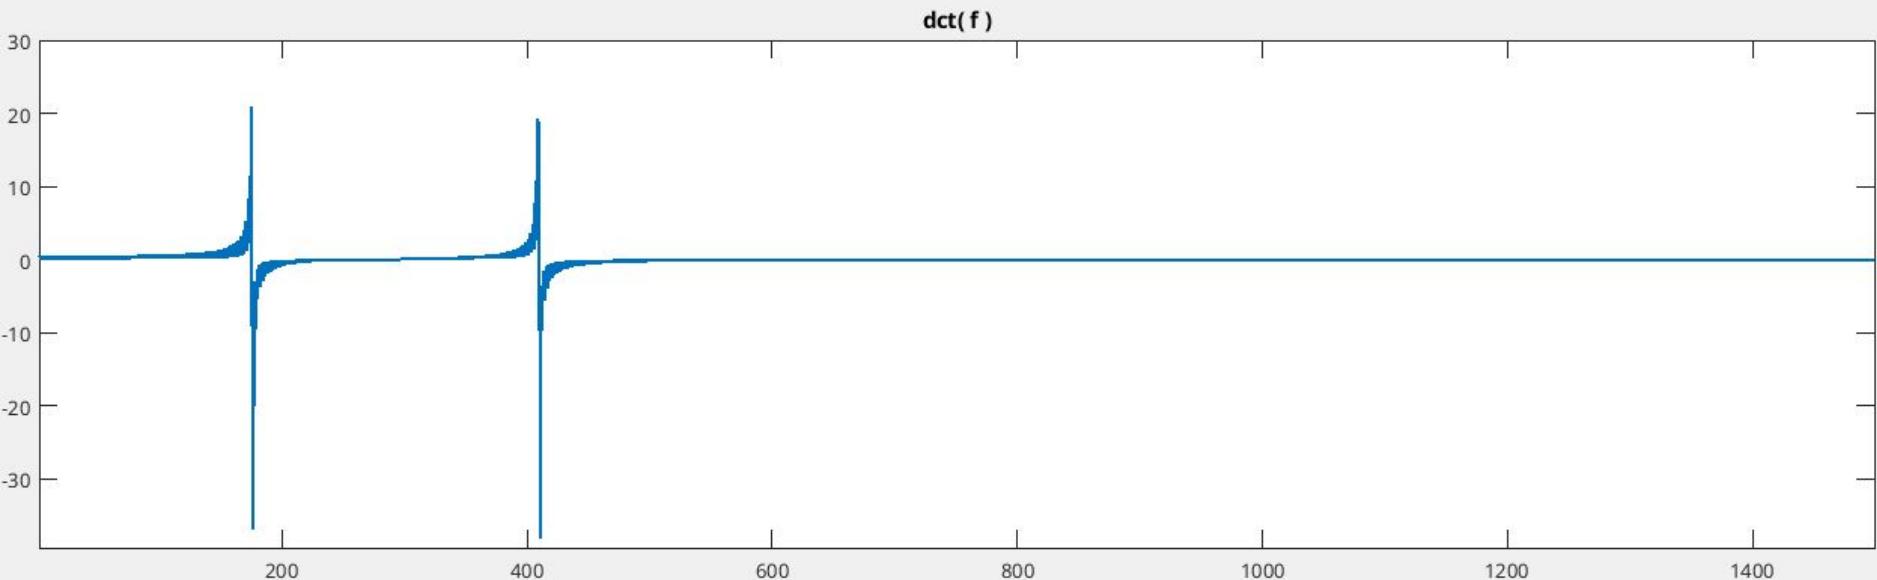
Comparing L¹- and L²-norm



Audio signal

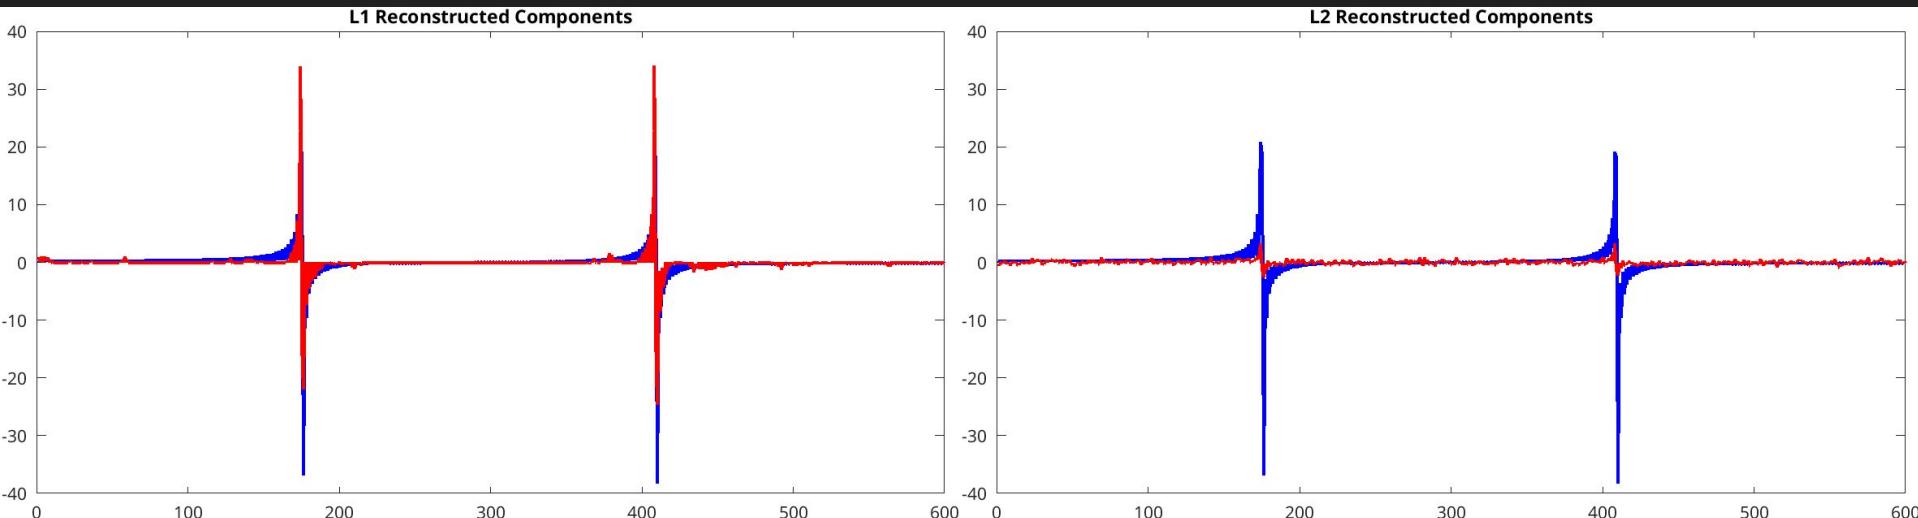
$$f = \sin(1394\pi t) + \sin(3266\pi t)$$

Comparing L¹- and L²-norm



DCT Components

Comparing L¹- and L²-norm



DCT Components

Comparing L₁- and L₂-norm

- L₂-norm is more sensitive to outliers, results in many small components
- L₁-norm does not punish outliers, results in more sparse results
 - We want sparsity!
- Solution: Constrain solution to minimum L₁ norm:

$$\text{minimize } \|\mathbf{x}\|_1$$

$$\text{s.t. } \mathbf{y} = \Phi \mathbf{x}$$

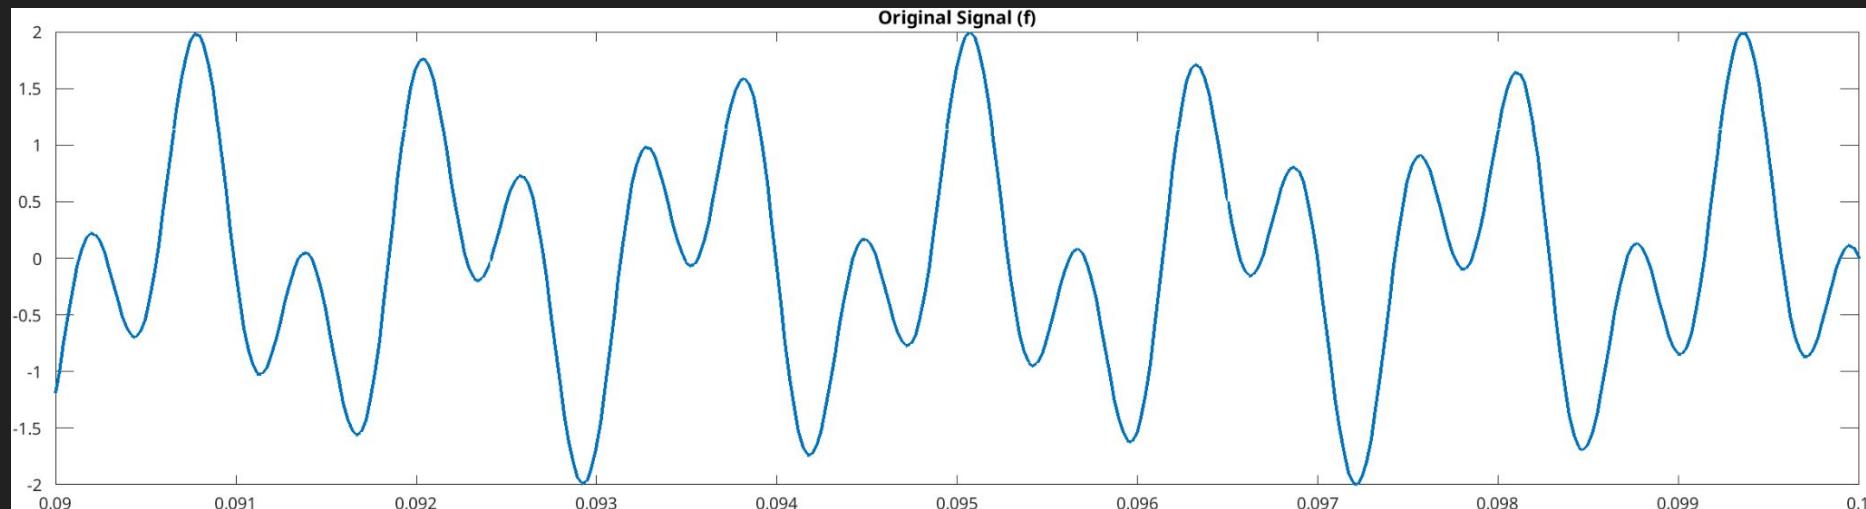
Calculating the Sensing Matrix (A or Φ)

- Start with $\mathbf{y} = \boldsymbol{\phi}\mathbf{f}$
 - \mathbf{f} : Target signal ($N \times 1$)
 - $\boldsymbol{\phi}$: Sampling/sparsifying matrix ($M \times N$)
 - \mathbf{y} : Compressed/undersampled signal ($M \times 1$)
- Let $\mathbf{f} = \boldsymbol{\psi}\mathbf{x}$
 - \mathbf{f} : Reconstructed signal ($N \times 1$)
 - $\boldsymbol{\psi}$: Transform matrix from spectral to temporal (frequency to time) domains ($N \times N$)
 - Inverse DCT of I_N
 - \mathbf{x} : Sparse frequency components ($N \times 1$)
 - Also the solution to the optimization problem
- Substitute \mathbf{f} : $\mathbf{y} = \boldsymbol{\phi}\boldsymbol{\psi}\mathbf{x}$
- From $\mathbf{y} = \mathbf{A}\mathbf{x} \rightarrow \mathbf{A} \equiv \boldsymbol{\phi}\boldsymbol{\psi}$
 - \mathbf{A} is just subsampled rows of the transform matrix!

1-Dimensional CS

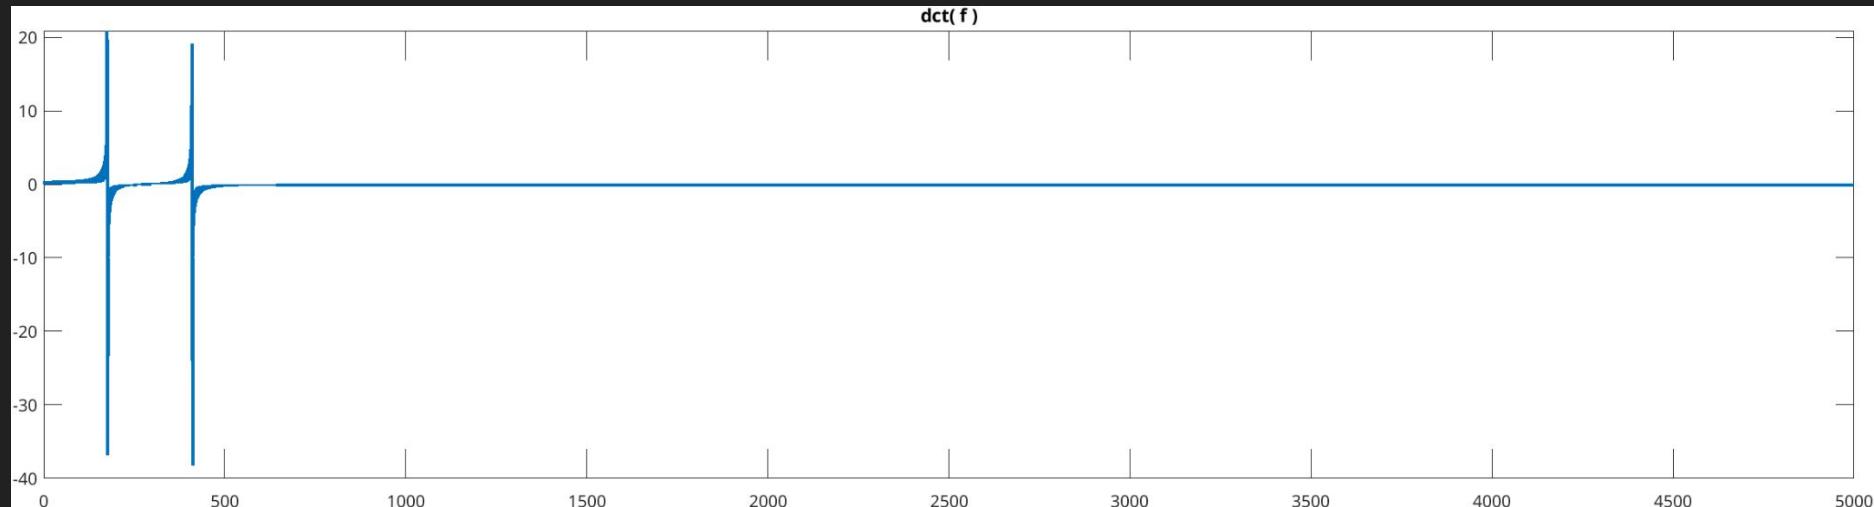
- Steps:
 1. Generate full signal of N samples
 2. Compute DCT components of original signal
 3. Extract M (e.g., $0.5N$) samples from original signal -- \mathbf{y}
 4. Generate sensing matrix \mathbf{A} (or Φ)
 5. Perform convex optimization
 6. Perform inverse DCT on optimized \mathbf{x} to retrieve reconstructed signal

Audio Signal Example



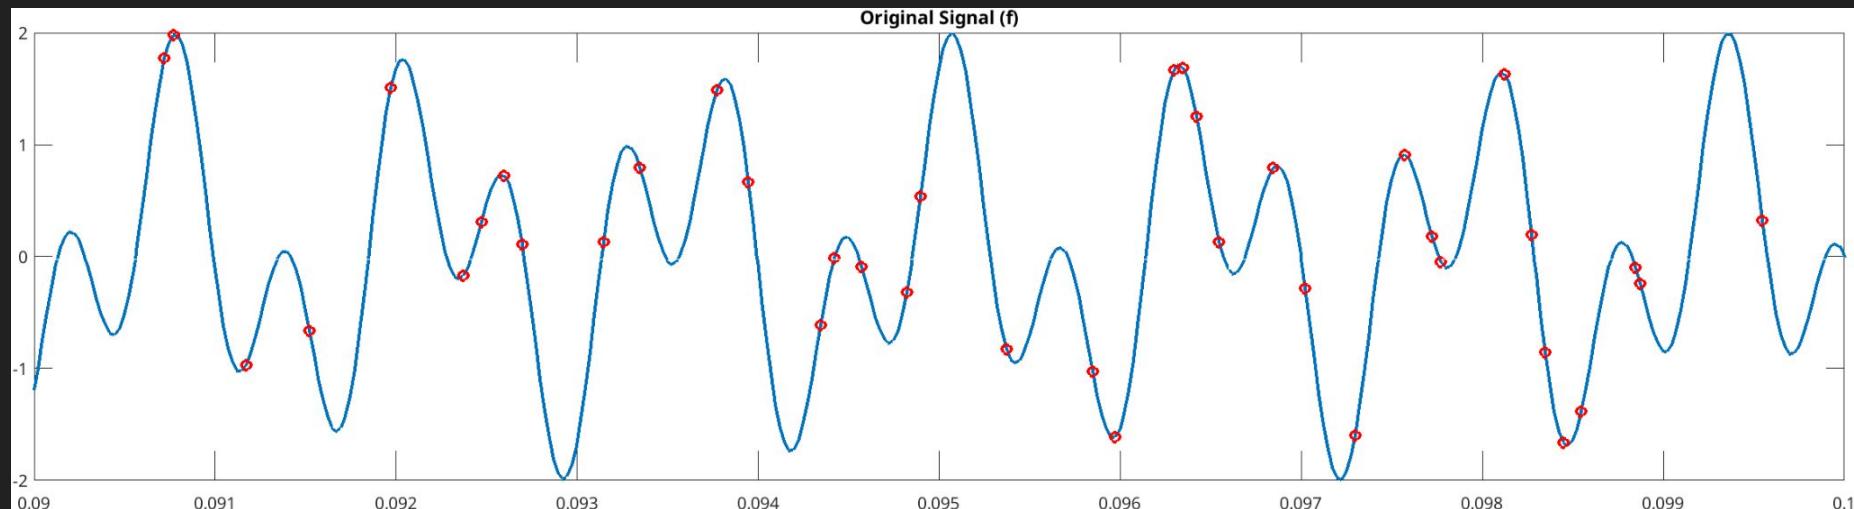
1. Generate full signal of $N = 5,000$ samples

Audio Signal Example



2. Compute DCT components of original signal

Audio Signal Example



3. Extract $M = 0.10N = 500$ samples from original signal

Audio Signal Example

4. Generate sampling matrix A (or Φ)

```
temp = randperm( n ); % Random permutation of [n]
ind = temp( 1:m ); % Grab the first 500 random indexes.

% Figure out mapping from time to frequency domain.
D = dct( eye( n, n ) );
A = D( ind, : );
```

```
cvx_begin
    variable x( n );
    minimize( norm( x, 1 ) );
    subject to
        A*x == fr.';
cvx_end
```

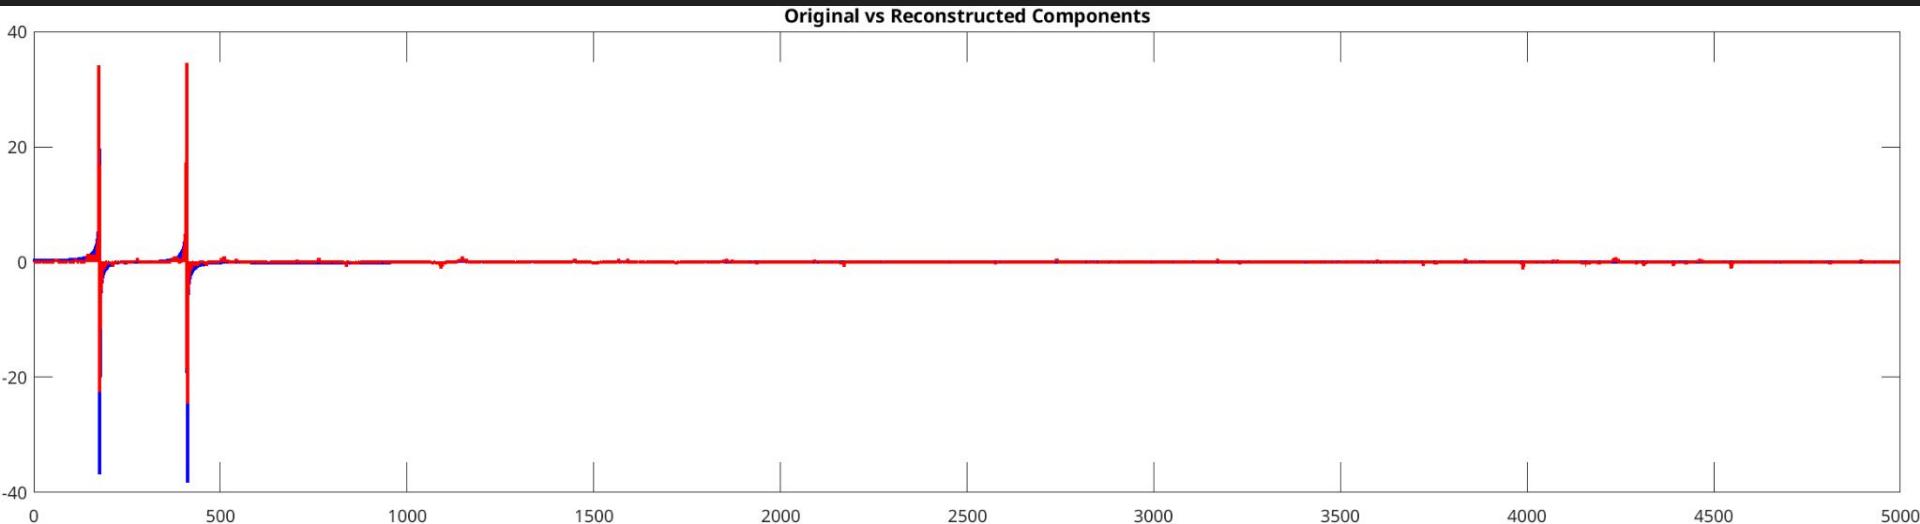
5. Perform convex optimization

6. Perform DCT on optimized x to retrieve reconstructed signal

```
sig1 = dct( x );
```

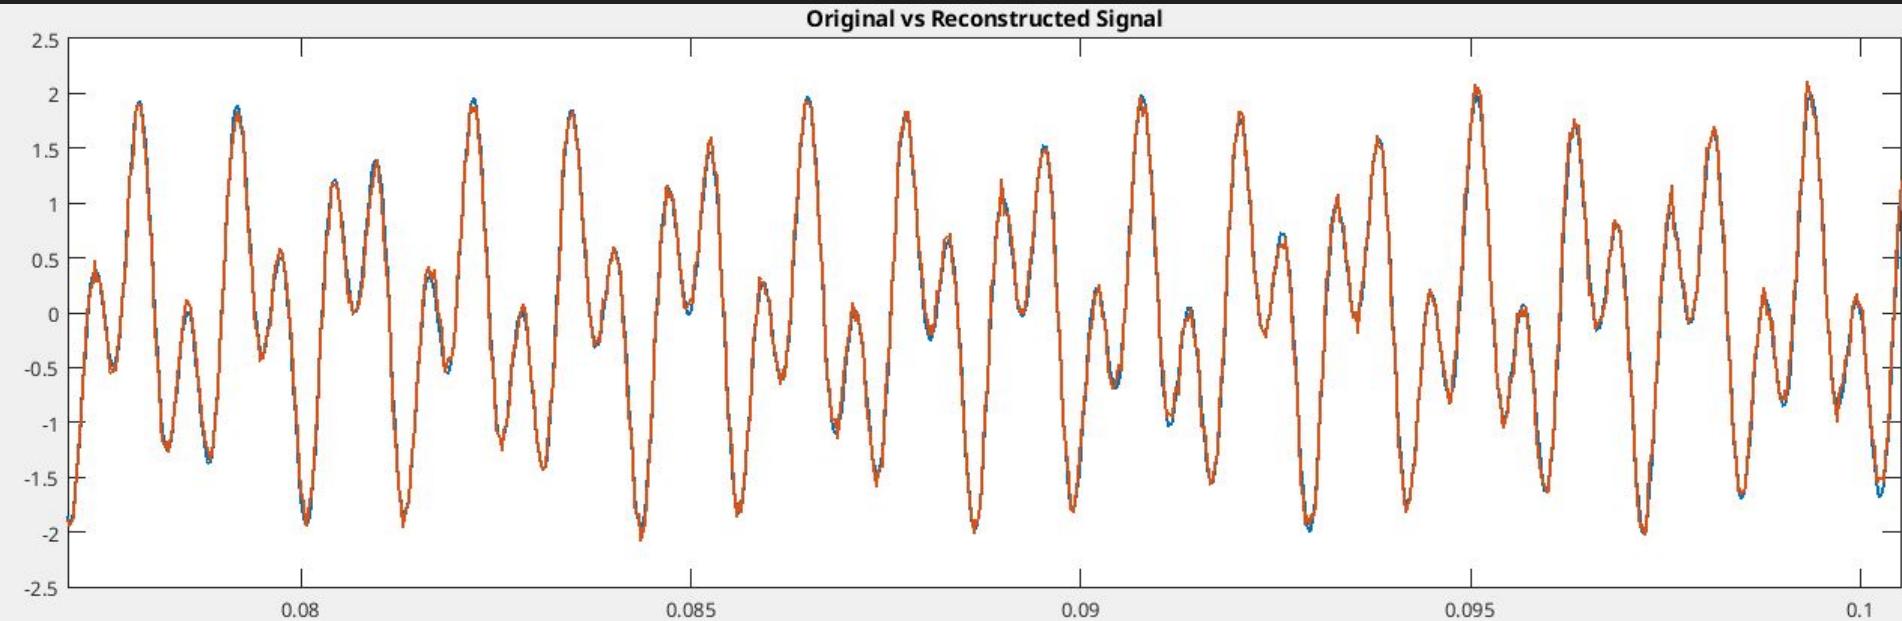
Audio Signal Example - Results

- Verify reconstructed DCT components



Audio Signal Example - Results

- Verify reconstructed signal



Application to 2D Signals (Images)

- CS in images requires flattening the images
 - Example: A 512x512 image becomes a 262,144x1 vector
 - Each color (RGB) is calculated separately (three vectors)
- Problem?
 - Requires massive memory!
 - Example: $N = 512 \times 512 = 262,144$, $M = 0.5N = 131,072$
 - \mathbf{A} is 262,144x131,072, \mathbf{x} is 262,144x1, \mathbf{y} is 131,072x1
 - That's 34.4 billion values
 - Assuming uint8, 34GB of RAM required (not including the requirements for the optimization problem!). What if we had color? 3 times that!
 - We can only do very small portions of an image at a time because of this.
- Solution?
 - L-BFGS and OWL-QN!

L-BFGS and OWL-QN

- Limited memory BFGS (Broyden-Fletcher-Goldfarb-Shanno) algorithm
 - Approximates BFGS with a focus on memory optimization
 - Minimizes function $f(\mathbf{x})$
- Orthant-Wise Limited-Memory Quasi-Newton (OWL-QN)
 - L-BFGS variant specific for fitting L_1 models using their sparsity properties
 - Developed by Microsoft Research
 - C library is available. Using numpy wrappers, can be included in Python
 - Minimizes the function $f(\mathbf{x}) = g(\mathbf{x}) + C \|\mathbf{x}\|_1$, where g is a differentiable convex loss function
 - Let's use the least squares objective function as our g :
 - $\|\mathbf{Ax} - \mathbf{y}\|_2^2$
 - $\nabla g = 2(\mathbf{A}^T \mathbf{Ax} - \mathbf{A}^T \mathbf{y})$

Original



50%

Mask



Reconstructed



30%



Comparison - Lena (512x512)

Original



10%

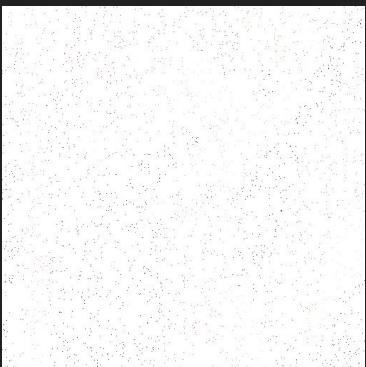
Mask



Reconstructed

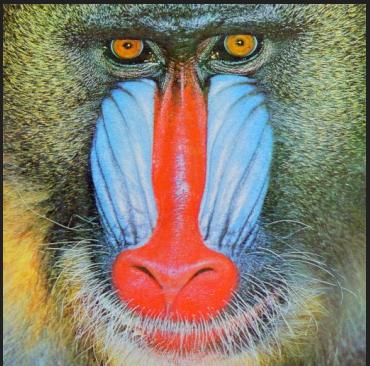


1%

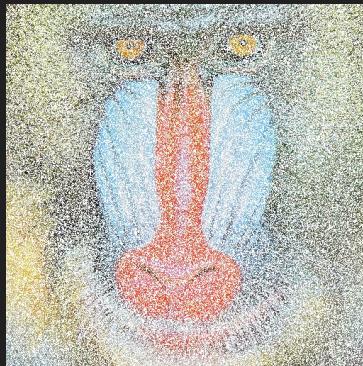


Comparison - Lena (512x512)

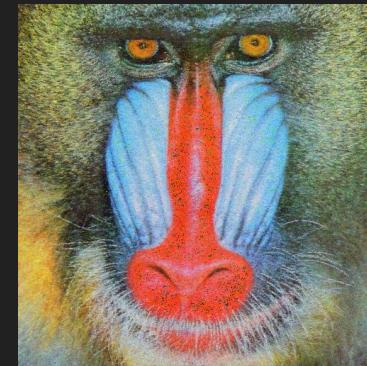
Original



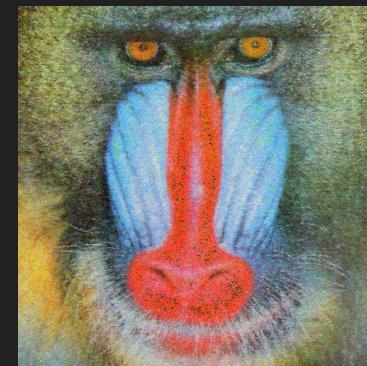
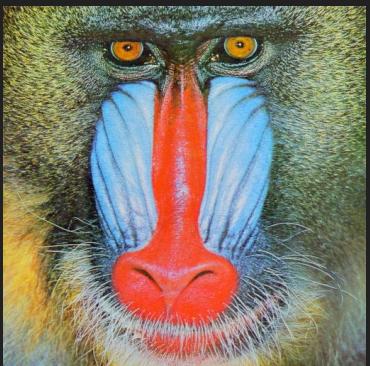
Mask



Reconstructed



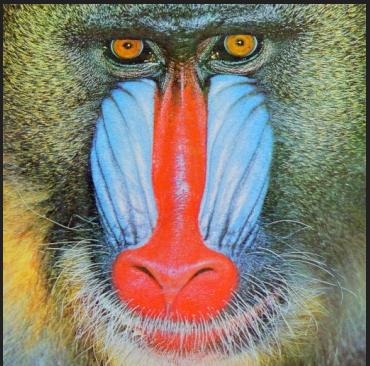
50%



30%

Comparison - Mandrill (512x512)

Original

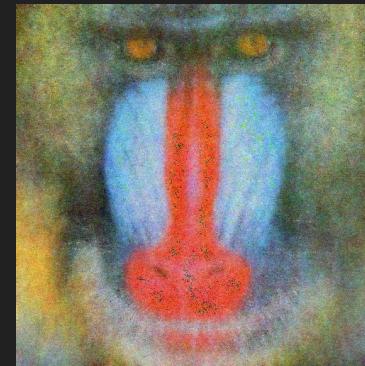


10%

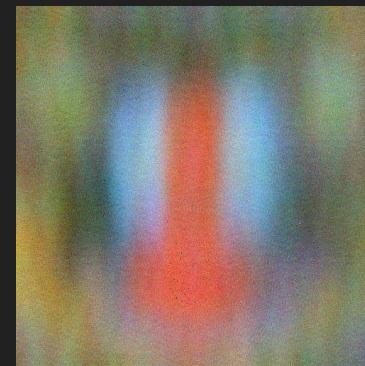
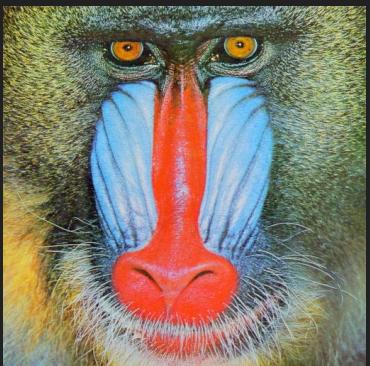
Mask



Reconstructed



1%



Comparison - Mandrill (512x512)

Original



Mask

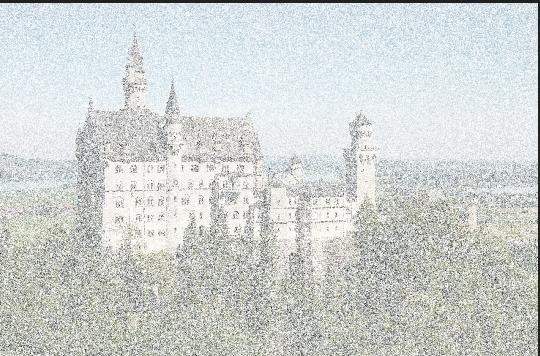


Reconstructed



50%

30%



Comparison - Neuschwanstein (1,280x838)

Original



Mask



Reconstructed



10%

1%



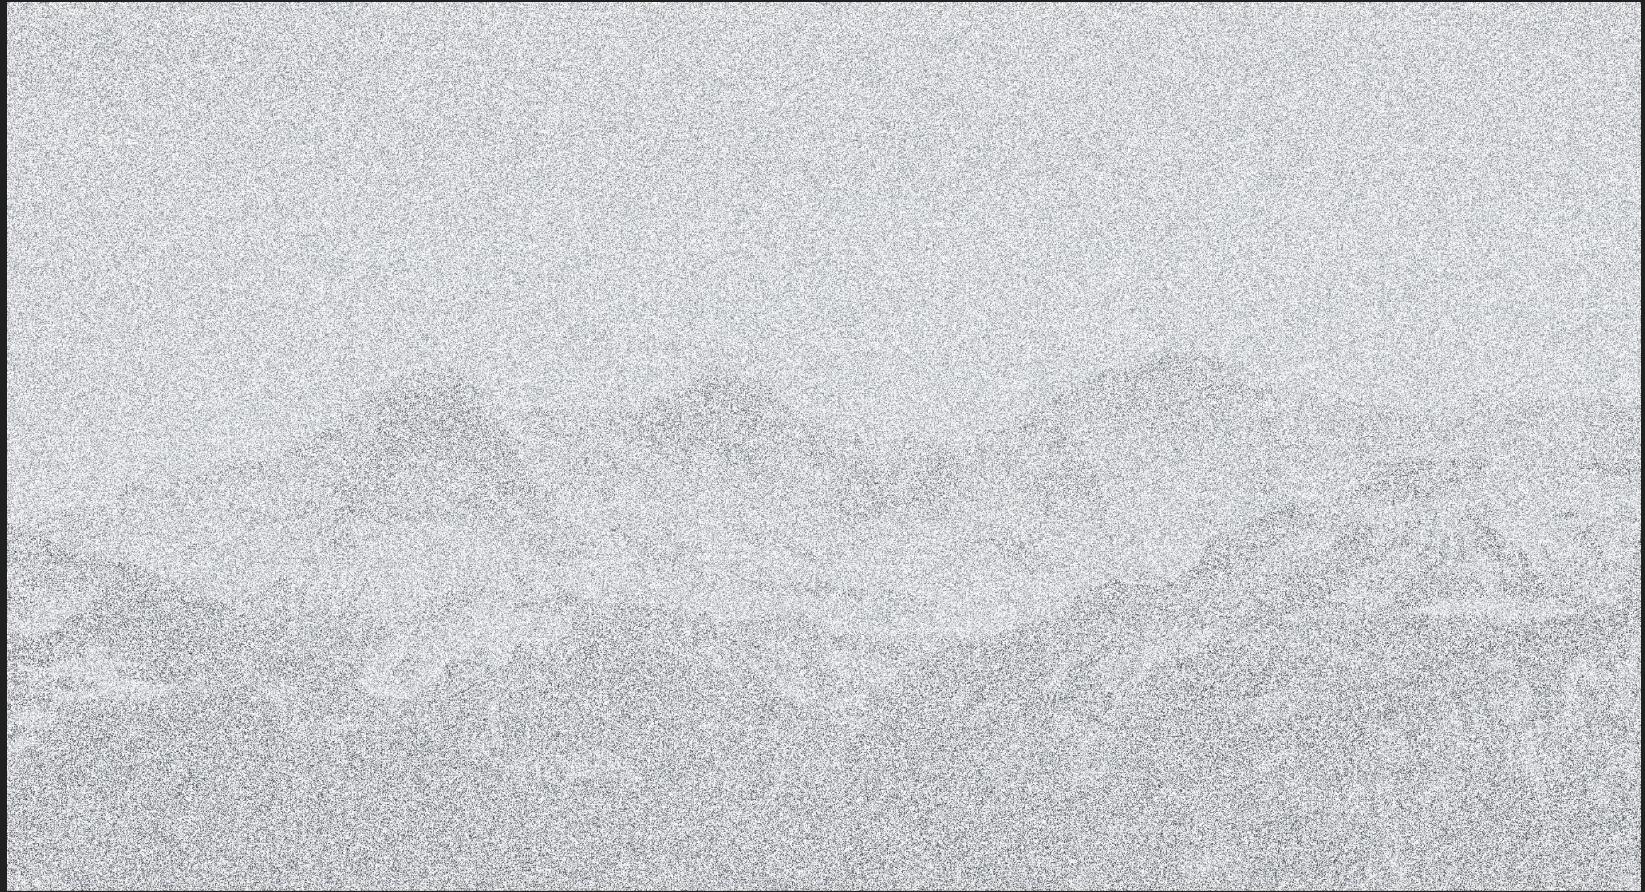
Comparison - Neuschwanstein (1,280x838)

Results

Sample Percentage	SSIM			MSE			PSNR		
	Lena	Mandrill	Neuschwanstein	Lena	Mandrill	Neuschwanstein	Lena	Mandrill	Neuschwanstein
100%	0.978	0.993	0.933	5.226	6.991	157.099	40.949	36.685	26.169
90%	0.959	0.932	0.770	10.243	74.015	376.106	38.026	29.438	22.378
80%	0.937	0.870	0.648	15.987	154.363	617.950	36.093	26.245	20.221
70%	0.911	0.800	0.554	24.849	255.127	871.792	34.178	24.063	18.727
60%	0.879	0.725	0.476	34.154	368.031	1147.691	32.796	22.472	17.533
50%	0.840	0.643	0.407	50.981	497.685	1426.921	31.057	21.161	16.587
40%	0.792	0.555	0.344	70.938	642.209	1728.594	29.622	20.054	15.754
30%	0.726	0.457	0.281	102.855	790.105	2021.594	28.009	19.154	15.074
20%	0.640	0.349	0.217	157.736	966.865	2319.439	26.151	18.277	14.477
10%	0.506	0.227	0.149	275.905	1163.247	2599.152	23.723	17.474	13.982
1%	0.251	0.105	0.093	1067.362	1685.916	2760.950	17.848	15.862	13.720



Original image (5,472x2,976)



30% sampled image



Recovered image

Original



Mask



30%

Comparison -- Mountains (5,472x2,976)

Original



Reconstructed



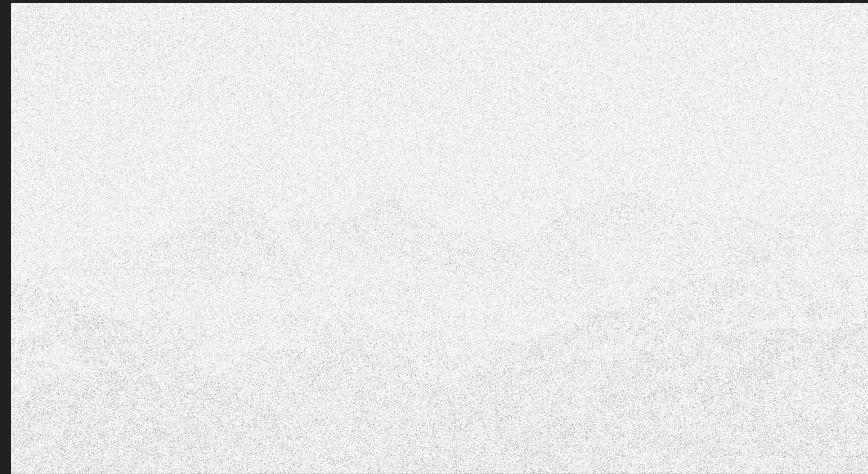
30%

Comparison -- Mountains (5,472x2,976)

Original



Mask



10%

Comparison -- Mountains (5,472x2,976)

Original



Reconstructed



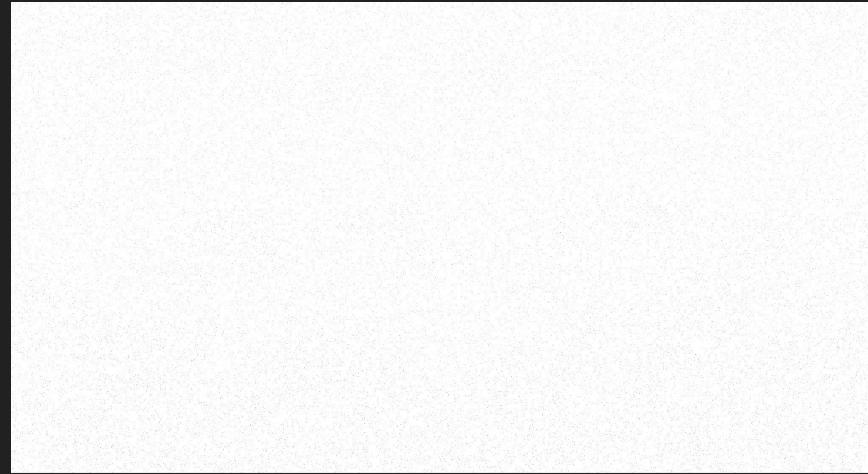
10%

Comparison -- Mountains (5,472x2,976)

Original



Mask



1%

Comparison -- Mountains (5,472x2,976)

1%

Original



Reconstructed



Comparison -- Mountains (5,472x2,976)

Applications

- Data compression
- Single pixel camera
- Medical imaging
- Wireless Sensor Networks
- Channel coding

Conclusion

- Compressive Sensing is a great tool for compression of data
 - Assumption is that data can be represented sparsely in some domain
- Algorithms need some optimization for efficient computation
 - Memory management
 - Long processing times
- Can greatly reduce the amount of time it takes to capture some data
 - Example: MRI

References

- [1] D. L. Donoho, "Compressed sensing," In *IEEE Transactions on Information Theory*, vol. 52, no. 4, pp. 1289-1306, April 2006. doi: 10.1109/TIT.2006.871582.
- [2] A. Hladnik, P. Saksida, "Compressed sensing and some image processing applications," In *Int. Symp. on Graphic Engineering and Design*, 2018, pp. 567-572. doi: 10.24867/GRID-2018-p68.
- [3] Nathan Kutz. *Compressive Sensing*. (Feb. 6, 2021). Accessed: Apr. 29, 2021. [Online Video]. Available: <https://www.youtube.com/watch?v=rt5mMEmZHfs>
- [4] P. C. Nahar, M. T. Kolte, "An introduction to compressive sensing and its applications," *Int. Jour. of Scientific and Research Publications*, vol. 4, no. 6. Jun. 2014. [Online]. Available: <http://www.ijrsp.org/research-paper-0614/ijrsp-p3076.pdf>
- [5] "Compressed Sensing in Python." Pyrunner. <http://www.pyrunner.com/weblog/2016/05/26/compressed-sensing-python/> (retrieved Apr. 29, 2021).
- [6] G. Andrew, J. Gao, "Scalable training of L1-regularized log-linear models," *Proceedings of the 24th International Conference on Machine Learning*, 2007, pp. 33-40, doi: 10.1145/1273496.1273501.
- [7] P. Boufounos, *Sparsity, randomness and compressed sensing* [PDF]. Available: <http://mlsp.cs.cmu.edu/courses/fall2009/class24/CS.slides.pdf>. (accessed Apr. 29, 2021).